



Instacart Market Basket Analysis

ML Portfolio Project

Maria Opekhtina

ML-B17, April 2023

Agenda

Company Intro and Background

Data

Visualizations

Code

Model Evaluation


Challenges and Next Steps

What is Instacart?

Instacart is an American delivery company that operates a grocery delivery and pick-up service in the United States and Canada. The company offers its services via a website and mobile app. The service allows customers to order groceries from participating retailers with the shopping being done by a personal shopper.

US Digital Ad Revenue Growth, by Company, 2023

% change



Note: includes advertising that appears on desktop and laptop computers as well as mobile phones, tablets, and other internet-connected devices, and includes all the various formats of advertising on those platforms; net ad revenues after companies pay traffic acquisition costs (TAC) to partner sites
Source: eMarketer, Nov 2022

- Instacart's success is largely due to its push into advertising, as well as growth for its Instacart+ membership program.
- US digital grocery sales grew by 15.8% in 2022 and are set to grow another 14.8% this year.
- New initiatives - Instacart Business and Instacart Health,
- In 2023, Instacart plans to roll out a chatbot powered by ChatGPT to answer food-related questions and help shoppers find products

How does it work?

✓ Pro

- available throughout the United States and Canada
- optional cost-saving Instacart+ subscription option
- same-day delivery available
- partners with several large grocery store chains

✗ Con


- not available in some rural areas
- without Instacart+, delivery fee and service fee apply
- Instacart prices may be more expensive than those in-store
- some may find the service difficult to use



Project Objective

To predict which products appeared in the user's latest order based on customer's purchase history

- Better recommendations
- Inventory Management
- Targeted advertisement
- Promotions
- Customer retention
- User satisfaction



instacart
Delivery & Curbside Pickup

use promo code:
COOPOAD10

Available Only at
Cascadilla St. Store

GET 10% OFF

*Offer expires 10/31/21. Instacart® and the Instacart carrot logo are trademarks of Maplebear Inc., d/b/a Instacart. One offer per account only. Delivery subject to availability. Minimum \$35 order. Discount will be applied to the total purchase price for all non-alcohol products, and excludes taxes, tips and/or fees. Not available in all zip/post codes.

Dataset

- The dataset is anonymized and contains a sample of over 3 million grocery orders from more than 200,000 Instacart users
- 6 tables with order and product information


	order_id	user_id	eval_set	order_number	order_dow	order_hour_of_day	days_since_prior_order	product_id	add_to_cart_order	reordered	product_name	aisle_id	department_id	aisle	department
0	2539329	1	prior	1	2	8	NaN	196.0	1.0	0.0	Soda	77	7	soft drinks	beverages
1	2539329	1	prior	1	2	8	NaN	14080.0	2.0	0.0	Organic Unsweetened Vanilla Almond Milk	91	16	soy lactosefree	dairy eggs
2	2539329	1	prior	1	2	8	NaN	12424.0	3.0	0.0	Original Beef Jerky	23	19	popcorn jerky	snacks
3	2539329	1	prior	1	2	8	NaN	26080.0	4.0	0.0	Aged White Cheddar Popcorn	23	19	popcorn jerky	snacks
4	2539329	1	prior	1	2	8	NaN	26400.0	5.0	0.0	XL Pick-A-Size Paper Towel Rolls	54	17	paper goods	household
...
20641986	2977660	206209	prior	13	1	12	7.0	14200.0	5.0	1.0	Tomato Paste	9	9	pasta sauce	dry goods pasta
20641987	2977660	206209	prior	13	1	12	7.0	38720.0	6.0	0.0	Brownie Crunch High Protein Bar	3	19	energy granola bars	snacks
20641988	2977660	206209	prior	13	1	12	7.0	31470.0	7.0	0.0	High Protein Bar Chunky Peanut Butter	3	19	energy granola bars	snacks
20641989	2977660	206209	prior	13	1	12	7.0	6570.0	8.0	0.0	Chocolate Peanut Butter Protein Bar	3	19	energy granola bars	snacks
20641990	2977660	206209	prior	13	1	12	7.0	22910.0	9.0	0.0	Roasted & Salted Shelled Pistachios	117	19	nuts seeds dried fruit	snacks

20641991 rows × 15 columns


Interesting Findings

- Most of the orders that were placed had been reordered items
- Most reordered items are fruits and vegetables

Reorders vs Non-reorders

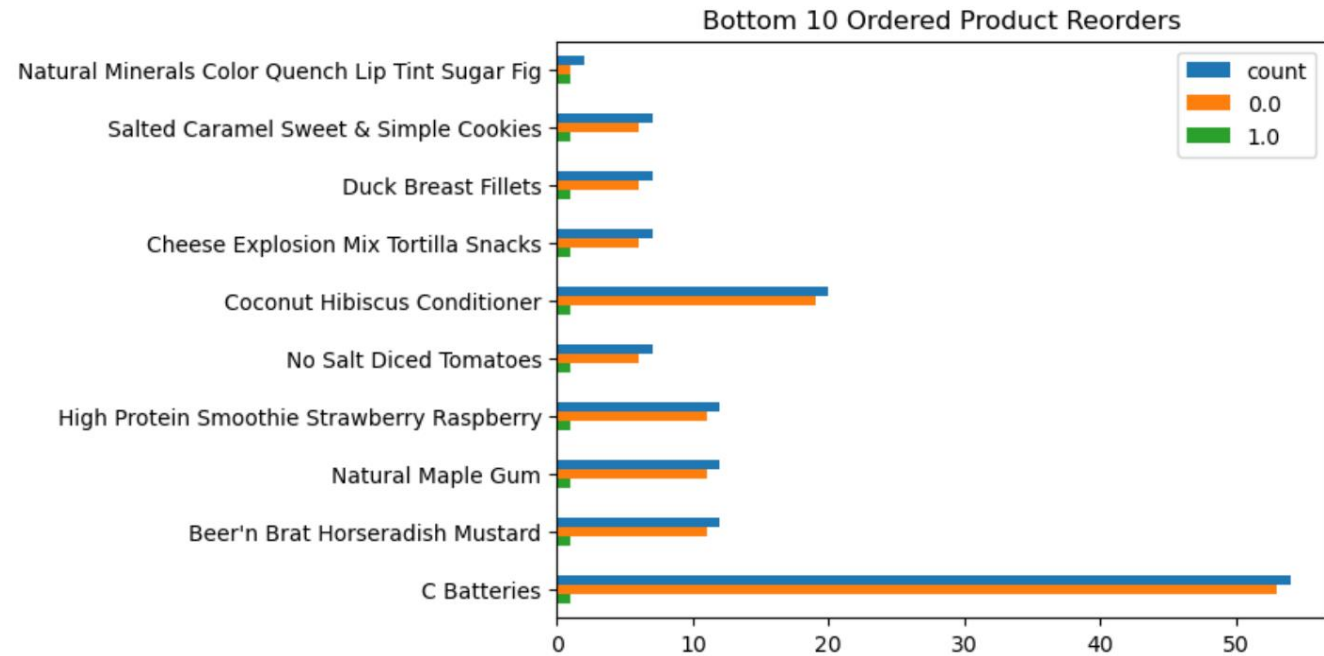
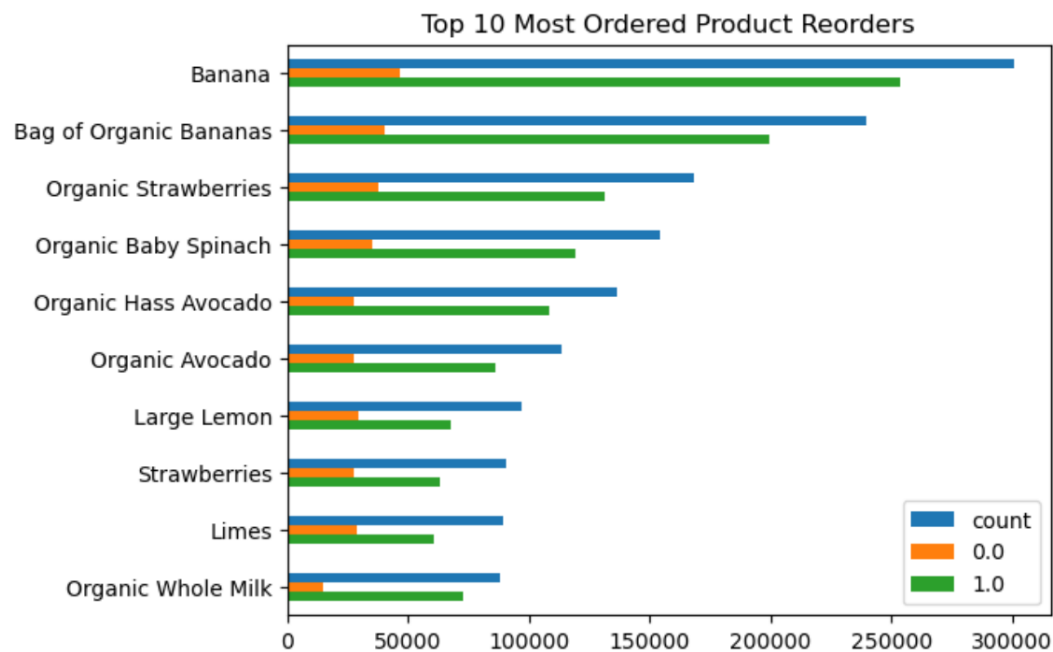


Top 10 Reordered Products

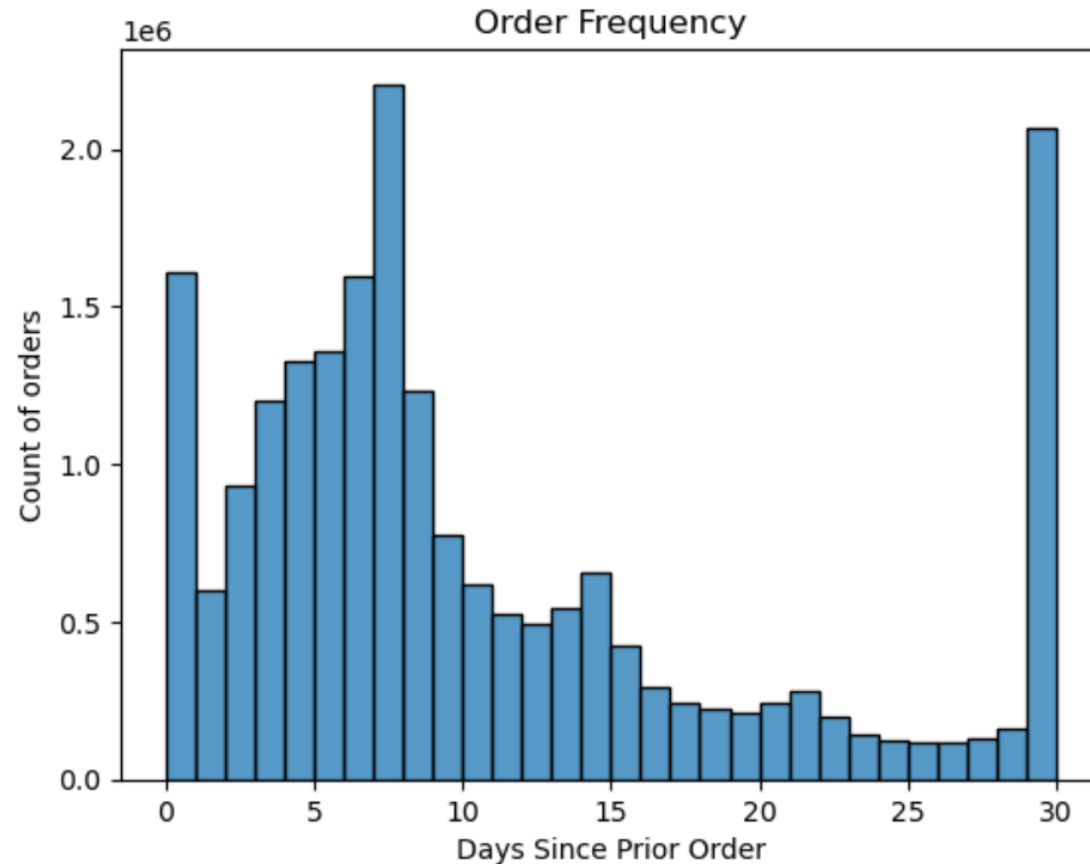


Most and Least Ordered Products

- Top 10 most ordered products are nicely aligned with top 10 reorders
- Bottom 10 are completely different, however indicating that people are likely not enjoying the items





Order Frequency



Peaks:

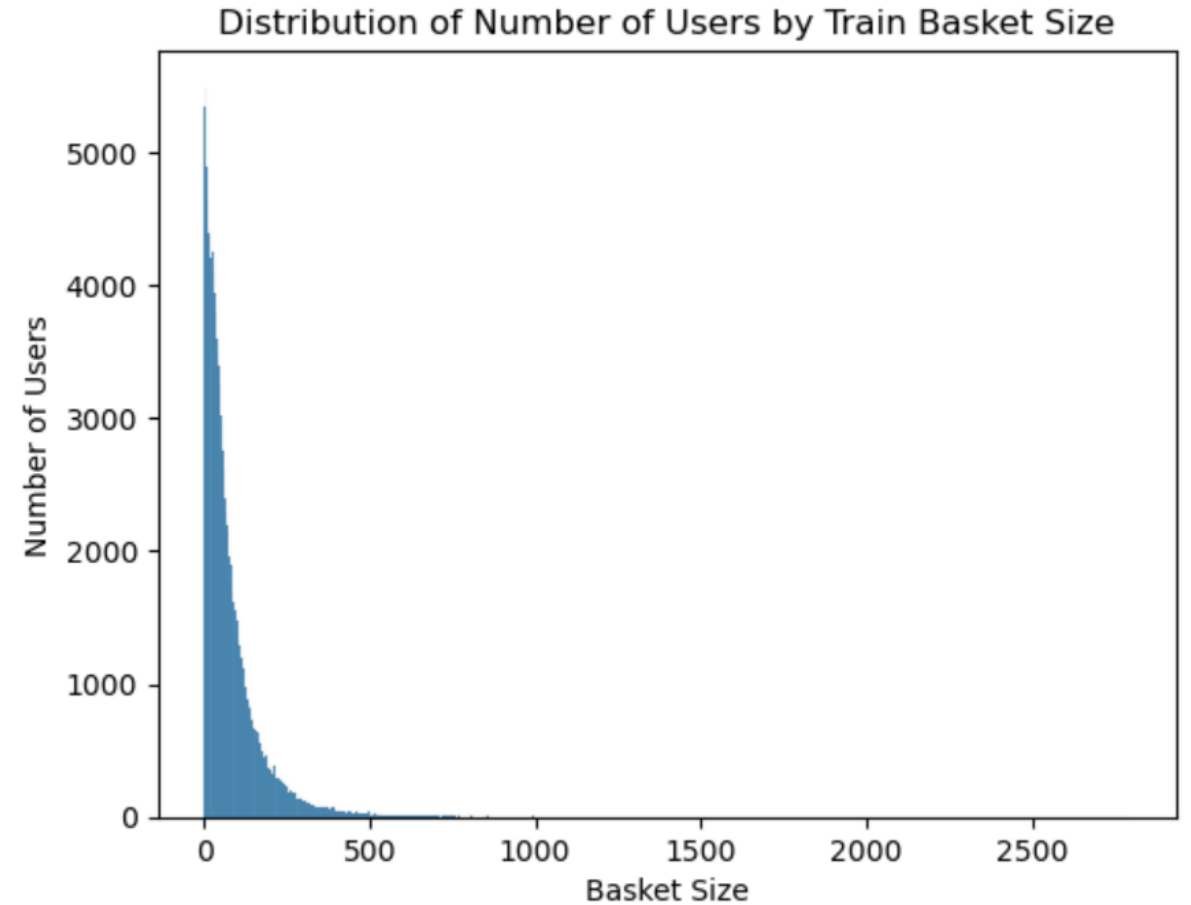

- 0 days since prior order:
 - indicates 1st order
- 7 days:
 - users shop weekly
- 30 days:
 - users shop monthly

When users shop



How Users Shop

Data is most likely being skewed by business accounts



Feature Engineering

Label

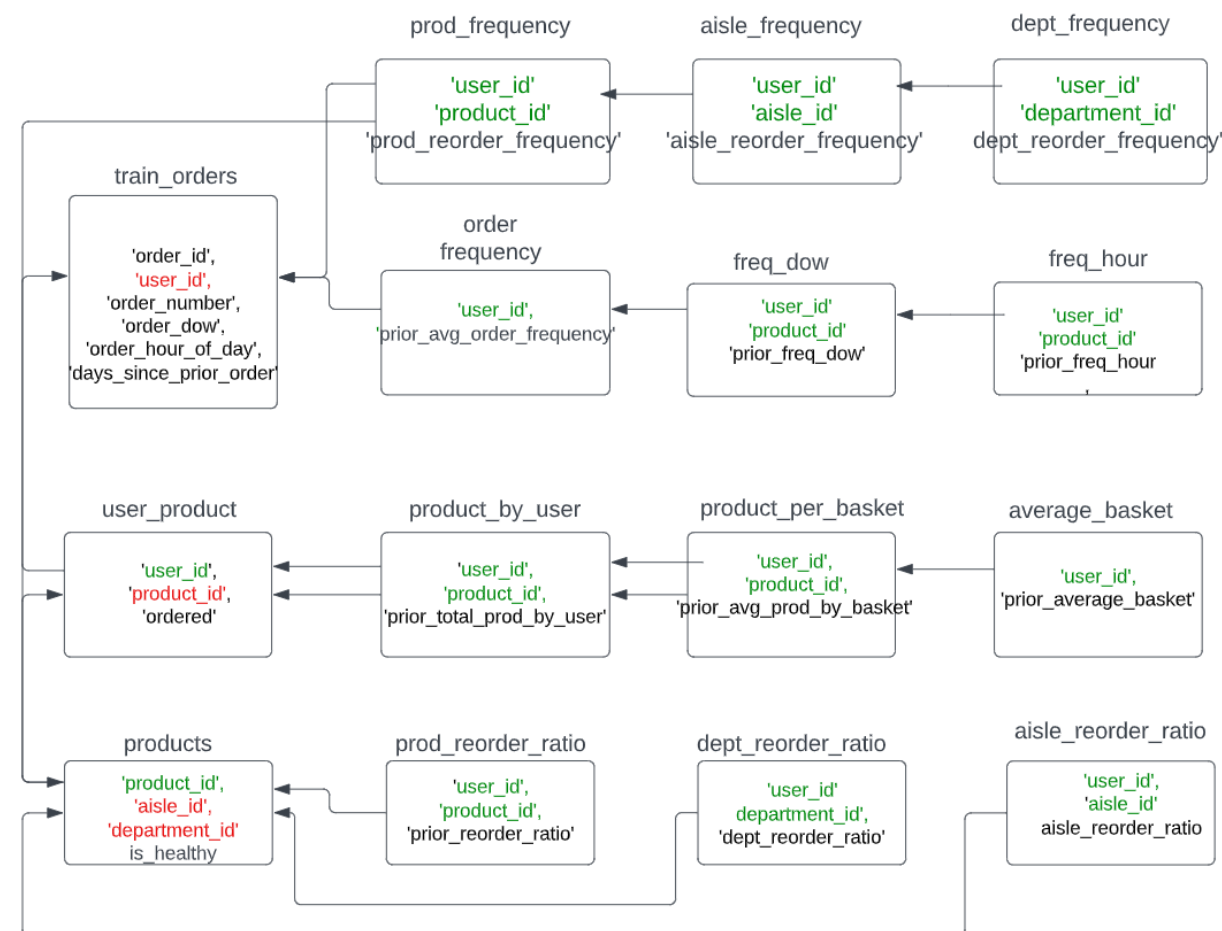
	user_id	product_id	ordered
0	1	196.0	1
1	1	14080.0	0
2	1	12424.0	0
3	1	26080.0	0
4	1	26400.0	0
...
8040479	206209	39232.0	0
8040480	206209	38720.0	0
8040481	206209	31472.0	0
8040482	206209	6568.0	0
8040483	206209	22912.0	0

8040484 rows × 3 columns

NLP

organic 5030
chocolate 2448
with 2236
free 2197
cheese 2095
chicken 1540
original 1452
sauce 1291
cream 1285
yogurt 1161
mix 1147
natural 1131
milk 1106
tea 1101
whole 1076

Features Schema



Final Table size:
1260057 rows × 23 columns

Models and Metrics


Models used:

- Logistic Regression
- Decision Tree
- XGBoost
- LightGBM
- Keras Model

Evaluation metrics:

- f1 score
 - target is very unbalanced, so relying on accuracy will not work since our model will predict 0 very precisely
- ROC AUC
 - a good indicator for whether the model is able to separate classes

Reorders vs Non-reorders



LightGBM

```
48]: %%time

model = lgb.LGBMClassifier()

pipe = Pipeline([
    ('preprocessing', preprocessor),
    ('rus', RandomUnderSampler(random_state=13)),
    ('model', model)
])

param_grid = {
    'model__random_state': [13],
    'model__max_depth': [5,10,15,20],
    'model__n_estimators': [40,140,240,340],
    'model__learning_rate': [0.5,.1,.2]
}

lgb_search = GridSearchCV(pipe, param_grid=param_grid, cv = 5, verbose = True, n_jobs = -1, scoring='f1_macro')
lgb_search.fit(X_train, y_train)

print(f'best parameters:', lgb_search.best_params_)
lgb_search.best_score_

Fitting 5 folds for each of 48 candidates, totalling 240 fits
best parameters: {'model__learning_rate': 0.2, 'model__max_depth': 5, 'model__n_estimators': 340, 'model__random_state': 13}
Wall time: 13min 40s

ut[548]: 0.5724775901723136
```

Logistic regression

```
%%time

model = LogisticRegression()

pipe = Pipeline([
    ('preprocessing', preprocessor),
    ('rus', RandomUnderSampler(random_state=13)),
    ('model', model)
])

param_grid = {
    'model__max_iter': [1000],
    'model__random_state': [13],
    'model__C': [.01,.05,.1,.5,1,5,10],
    'model__penalty': ['l2']
}

lr_search = GridSearchCV(pipe, param_grid=param_grid, cv = 5, verbose = True, n_jobs = -1, scoring='f1_macro')

lr_search.fit(X_train, y_train)

print(f'best parameters:', lr_search.best_params_)
lr_search.best_score_

Fitting 5 folds for each of 7 candidates, totalling 35 fits
best parameters: {'model__C': 0.01, 'model__max_iter': 1000, 'model__penalty': 'l2', 'model__random_state': 13}
Wall time: 48.3 s

0]: 0.3820393219623558
```

Decision Tree

```
%%time

model = DecisionTreeClassifier()

pipe = Pipeline([
    ('preprocessing', preprocessor),
    ('rus', RandomUnderSampler(random_state=13)),
    ('model', model)
])

param_grid = {
    'model__random_state': [13],
    'model__max_depth': [5,10,15,20],
    'model__min_samples_split': [2,3,4,5]
}

dt_search = GridSearchCV(pipe, param_grid=param_grid, cv = 5, verbose = True, n_jobs = -1, scoring='f1_macro')
dt_search.fit(X_train, y_train)

print(f'best parameters:', dt_search.best_params_)
dt_search.best_score_

Fitting 5 folds for each of 16 candidates, totalling 80 fits
best parameters: {'model__max_depth': 10, 'model__min_samples_split': 3, 'model__random_state': 13}
Wall time: 1min 51s
```

XGBoost

```
%%time

model = XGBClassifier()

pipe = Pipeline([
    ('preprocessing', preprocessor),
    ('rus', RandomUnderSampler(random_state=13)),
    ('model', model)
])

param_grid = {
    'model__random_state': [13],
    'model__max_depth': [5,10,15,20],
    'model__n_estimators': [40,140,240,340],
    'model__eta': [.05,.1,.2]
}

xgb_search = GridSearchCV(pipe, param_grid=param_grid, cv = 5, verbose = True, n_jobs = -1, scoring='f1_macro')
xgb_search.fit(X_train, y_train)

print(f'best parameters:', xgb_search.best_params_)
xgb_search.best_score_

Fitting 5 folds for each of 48 candidates, totalling 240 fits
best parameters: {'model__eta': 0.2, 'model__max_depth': 5, 'model__n_estimators': 340, 'model__random_state': 13}
Wall time: 53min 7s

4]: 0.5729772010981871
```

Pipeline and Stacking

Pipeline

Steps:

- ColumnTransformer:
 - separated columns into nums = numeric, one_hot_cats = categories to one-hot encode and freq_cats = categories to frequency encode
 - MinMaxScaler() performed better during test runs compared to standard scaler
 - There are no huge outliers in this dataset due to it's nature, so scaling should be sufficient
 - RareLabelEncoder() needed to be used on freq_cats before CountEncoder due to the CountEncoder() creating NaN values
- RandomUnderSampler()
 - performed better during test runs compared to over sampling
- Model - variable to be reset

```
nums = ['order_number', 'days_since_prior_order', 'order_size', 'prior_total_prod_by_user',
        'prior_avg_prod_by_basket', 'prior_average_basket', 'prior_avg_order_frequency',
        'product_reorder_ratio', 'dept_reorder_ratio', 'aisle_reorder_ratio',
        'prod_reorder_frequency', 'aisle_reorder_frequency', 'dept_reorder_frequency']

one_hot_cats = ['order_dow', 'order_hour_of_day', 'prior_freq_dow', 'prior_freq_hour', 'is_healthy']

freq_cats = ['user_id', 'product_id', 'aisle_id', 'department_id']
```

```
transformers = [('scaler', MinMaxScaler(), nums),
                ('oh_encoding', OneHotEncoder(drop='first'), one_hot_cats),
                ('rare_encoding', RareLabelEncoder(replace_with=0), freq_cats),
                ('freq_encoding', CountEncoder(), freq_cats)]

preprocessor = ColumnTransformer(transformers=transformers, remainder='drop')
```

```
pipe = Pipeline([
    ('preprocessing', preprocessor),
    ('rus', RandomUnderSampler(random_state=13)),
    ('model', model)
])
```

Stacking

Previously fitted model with the best parameters were stacked to check if they would improve the overall performance.

Unfortunately, there was not improvement observed neither in f1 scoring, nor the AUC. The model is unable to effectively differentiate between the classes, no matter which final estimator was used


```
estimators = [
    ('lr', lr.named_steps['model']),
    ('dt', dt.named_steps['model']),
    ('lgb', lgb_model.named_steps['model']),
    ('xgb', xgb.named_steps['model'])
]

stacked_pipe = Pipeline([
    ('preprocessing', preprocessor),
    ('rus', RandomUnderSampler(random_state=13)),
    ('clf', StackingClassifier(estimators=estimators, final_estimator=lgb.LGBMClassifier()))
])

stacked_pipe.fit(X_train, y_train)

y_pred_val = stacked_pipe.predict(X_val)

reports(y_pred_val, y_val)
```



Neural Networks – MLP Sequential

Model Training and validation

```
np.random.seed(13)
tf.random.set_seed(13)

mlp = Sequential()

mlp.add(InputLayer(input_shape=(76, )))
mlp.add(Dense(300, activation='relu'))
mlp.add(Dense(150, activation='relu'))

mlp.add(Dense(2, activation='sigmoid'))
```

```
mlp.summary()
```


Model: "sequential_75"

Layer (type)	Output Shape	Param #
=====		
dense_59 (Dense)	(None, 300)	23100
dense_60 (Dense)	(None, 150)	45150
dense_61 (Dense)	(None, 2)	302
=====		
Total params: 68,552		
Trainable params: 68,552		
Non-trainable params: 0		
=====		

```
mlp.compile(loss='sparse_categorical_crossentropy', optimizer='sgd', metrics=[get_f1])
```


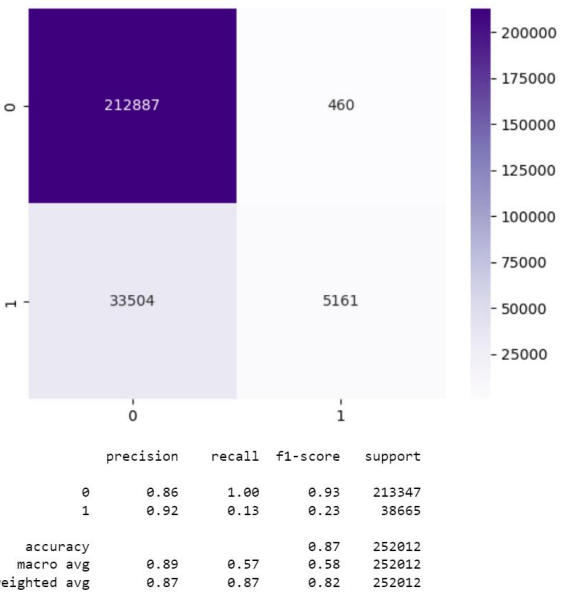
```
history = mlp.fit(X_resampled, y_resampled, epochs=20, validation_data=(X_val_scaled, y_val))

Epoch 1/20
1054/1054 [=====] - 28s 26ms/step - loss: 0.5347 - get_f1: 0.6301 - val_loss: 0.6300 - val_get_f1: 0.0436
Epoch 2/20
1054/1054 [=====] - 29s 28ms/step - loss: 0.4376 - get_f1: 0.6432 - val_loss: 0.4663 - val_get_f1: 0.0418
Epoch 3/20
1054/1054 [=====] - 29s 28ms/step - loss: 0.4229 - get_f1: 0.6462 - val_loss: 0.4063 - val_get_f1: 0.0414
Epoch 4/20
1054/1054 [=====] - 29s 27ms/step - loss: 0.4145 - get_f1: 0.6513 - val_loss: 0.3730 - val_get_f1: 0.0413
Epoch 5/20
1054/1054 [=====] - 29s 27ms/step - loss: 0.4076 - get_f1: 0.6553 - val_loss: 0.3501 - val_get_f1: 0.0413
Epoch 6/20
1054/1054 [=====] - 29s 28ms/step - loss: 0.4011 - get_f1: 0.6578 - val_loss: 0.3361 - val_get_f1: 0.0414
Epoch 7/20
1054/1054 [=====] - 29s 28ms/step - loss: 0.3941 - get_f1: 0.6610 - val_loss: 0.3335 - val_get_f1: 0.0418
Epoch 8/20
1054/1054 [=====] - 29s 27ms/step - loss: 0.3869 - get_f1: 0.6613 - val_loss: 0.3440 - val_get_f1: 0.0420
Epoch 9/20
1054/1054 [=====] - 30s 29ms/step - loss: 0.3797 - get_f1: 0.6627 - val_loss: 0.3591 - val_get_f1: 0.0421
Epoch 10/20
1054/1054 [=====] - 30s 29ms/step - loss: 0.3736 - get_f1: 0.6629 - val_loss: 0.3406 - val_get_f1: 0.0421
Epoch 11/20
1054/1054 [=====] - 29s 28ms/step - loss: 0.3677 - get_f1: 0.6625 - val_loss: 0.3416 - val_get_f1: 0.0421
Epoch 12/20
1054/1054 [=====] - 29s 27ms/step - loss: 0.3630 - get_f1: 0.6622 - val_loss: 0.3464 - val_get_f1: 0.0420
Epoch 13/20
```

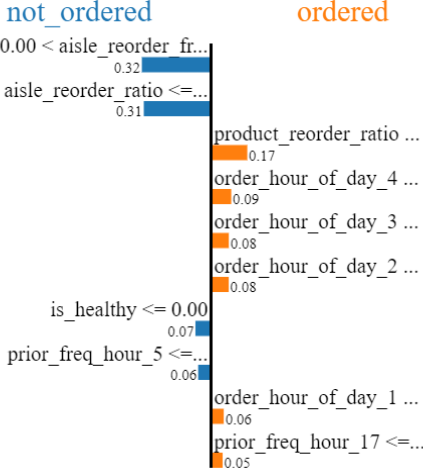
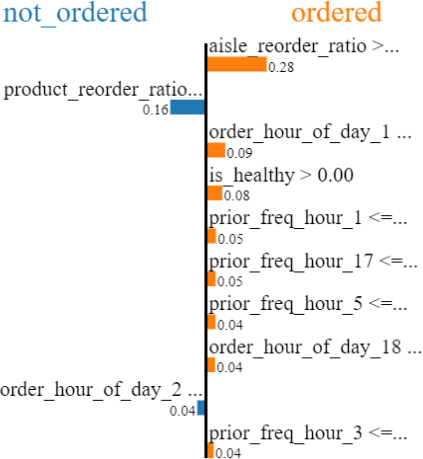


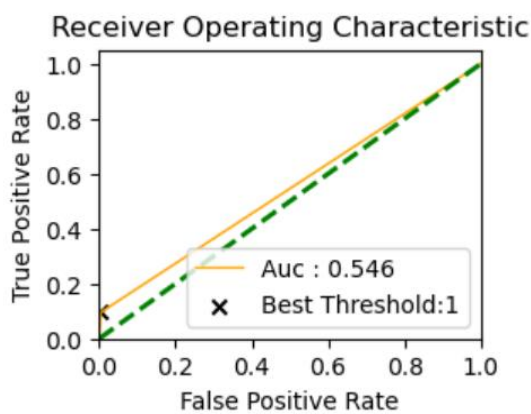
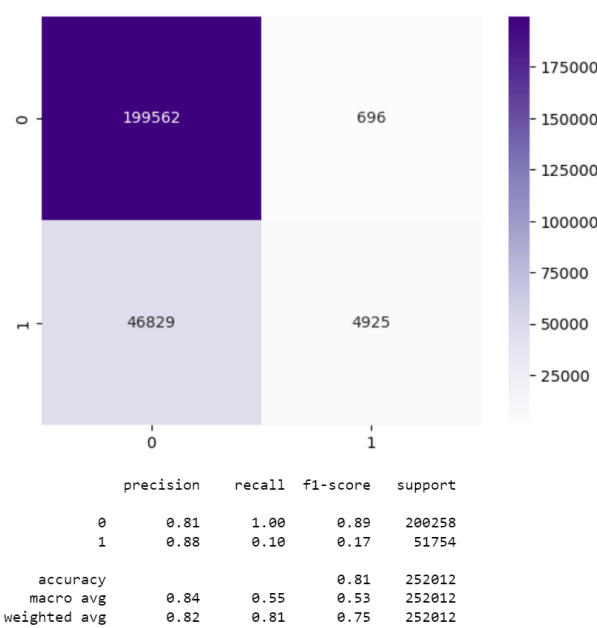


Model Evaluation

Stacked Model



Keras Model



Feature	Value
aisle_reorder_ratio	0.76
product_reorder_ratio	0.93
order_hour_of_day_1	0.00
is_healthy	1.00
prior_freq_hour_1	0.00
prior_freq_hour_17	0.00
prior_freq_hour_5	0.00
order_hour_of_day_18	0.00
order_hour_of_day_2	0.00

Feature	Value
aisle_reorder_frequency	2.00
aisle_reorder_ratio	0.00
product_reorder_ratio	0.00
order_hour_of_day_4	0.00
order_hour_of_day_3	0.00
order_hour_of_day_2	0.00
is_healthy	0.00
prior_freq_hour_5	0.00
order_hour_of_day_1	0.00

Challenges

- Unfamiliar problem
- A large chunk of time was spent on data preparation and analysis
- Highly imbalanced classes
- Difficult to interpret findings when using blackbox models
- Challenging to navigate a pipeline
- Models did not end up performing as well as intended and the scores were way below the average of 0.35 f1 score.

Next Steps

Some things that can improve the model performance:

- adding more/better features will definitely help
- trying different types of encoders
- trying other ML algorithms that people have used: RandomForest, Catboost
- Try a wider scope in terms of hyperparameter tuning, trying wider range of parameters and re-running the SearchGrid to optimize
- Association rules search with Apriori Algorithm